

High Performance Computing

LECTURE #3



Agenda

- Parallel Computing Platform (Logical Organization)
- Flynn Taxonomy



1- Introduction

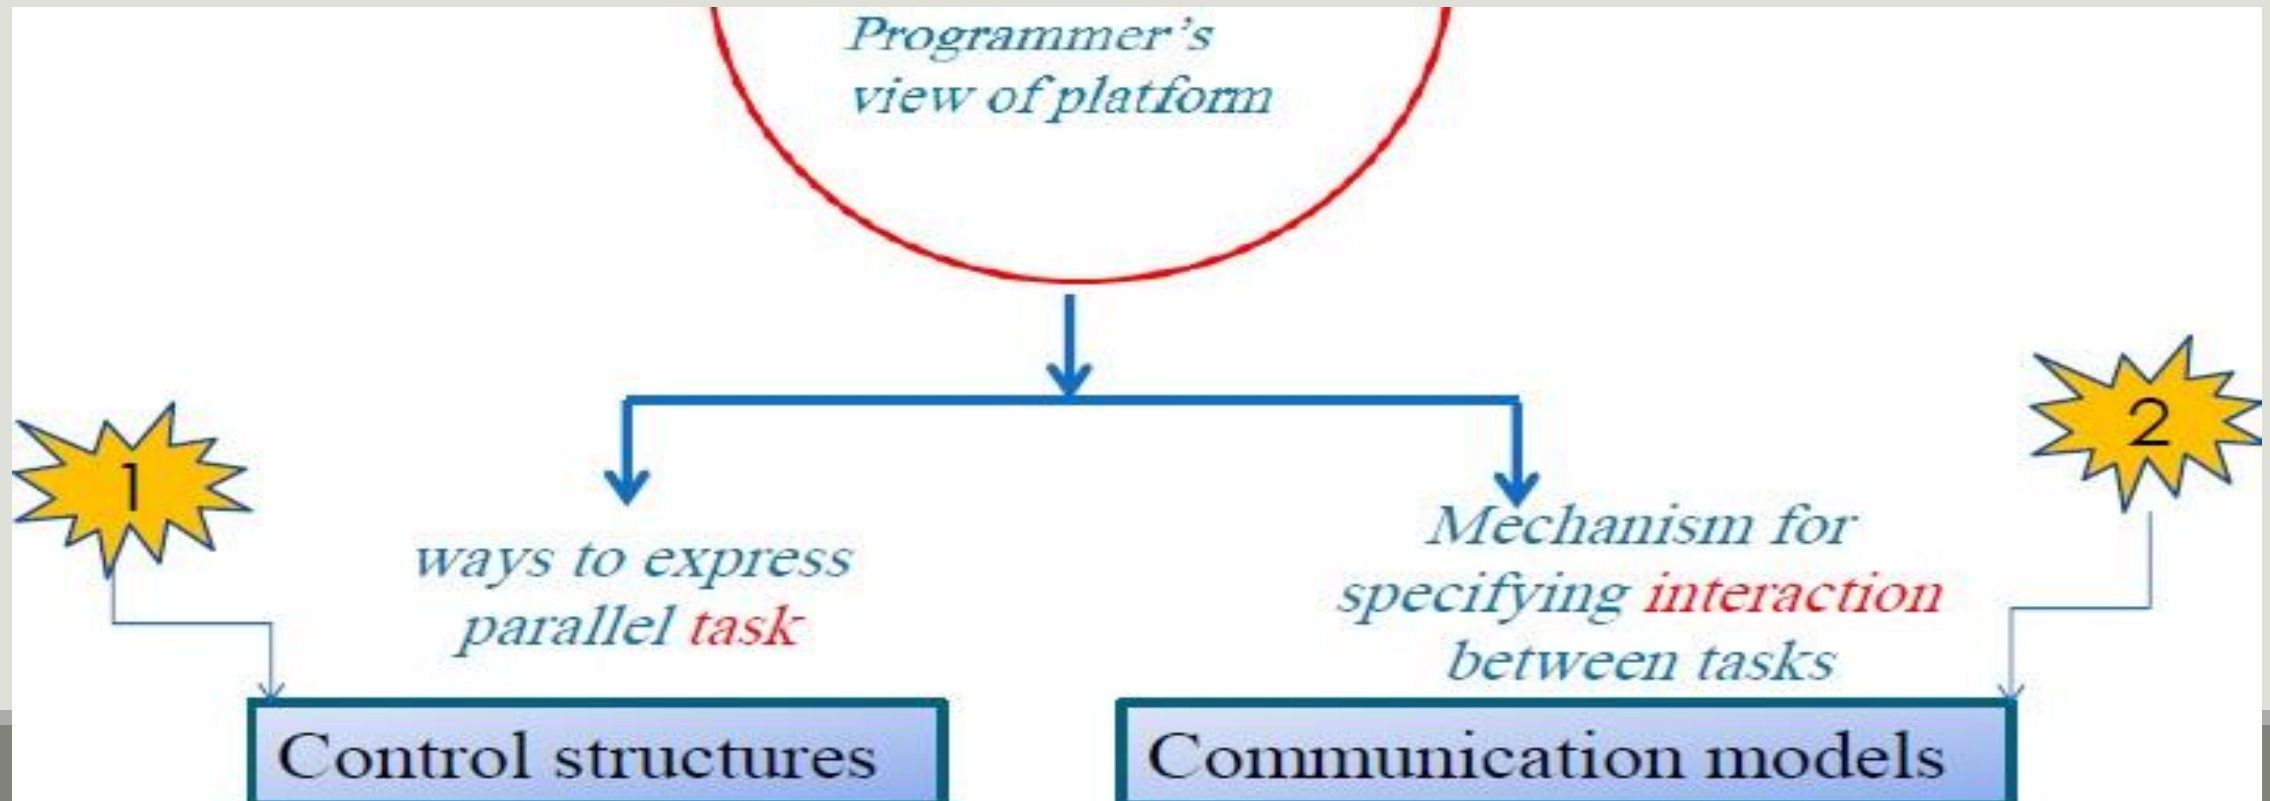
- ❖ A **computing platform** includes a **hardware architecture** and a **software framework** (including application frameworks), where the combination allows software to run.
- ❖ Typical platforms include a
 - computer architecture,
 - operating system,
 - programming languages and
 - related user interface (run-time system libraries or graphical user interface).

Parallel Computing Platforms

Logical organization is **split** into **exactly two non-overlapping parts**

Control structures Communication models

An explicitly parallel program must specify accurately the interaction between concurrent subtasks..



❖ Parallelism can be expressed at various levels of **granularity**.

❖ **Computation / Communication Ratio:**



- In parallel computing, granularity is a **qualitative** measure of the ratio of computation to communication.
- Periods of computation are typically separated from periods of communication by **synchronization events**.



Fine-grain Parallelism:

- ❖ Relatively small amounts of computational work are done between communication events
- ❖ Low computation to communication ratio.
- ❖ Facilitates **load balancing**.
- ❖ Implies high **communication overhead**:

If granularity is too fine it is possible that the **overhead** required for **communications** and **synchronization** between tasks takes longer than the computation.




Coarse-grain Parallelism:

- ❖ Relatively large amounts of computational work are done between communication/synchronization events
- ❖ High computation to communication ratio
- ❖ **Harder to load balance** efficiently.




Fine-grain Parallelism:

❖ Example: Dense matrix multiplication



Coarse-grain Parallelism:



Which is Best?

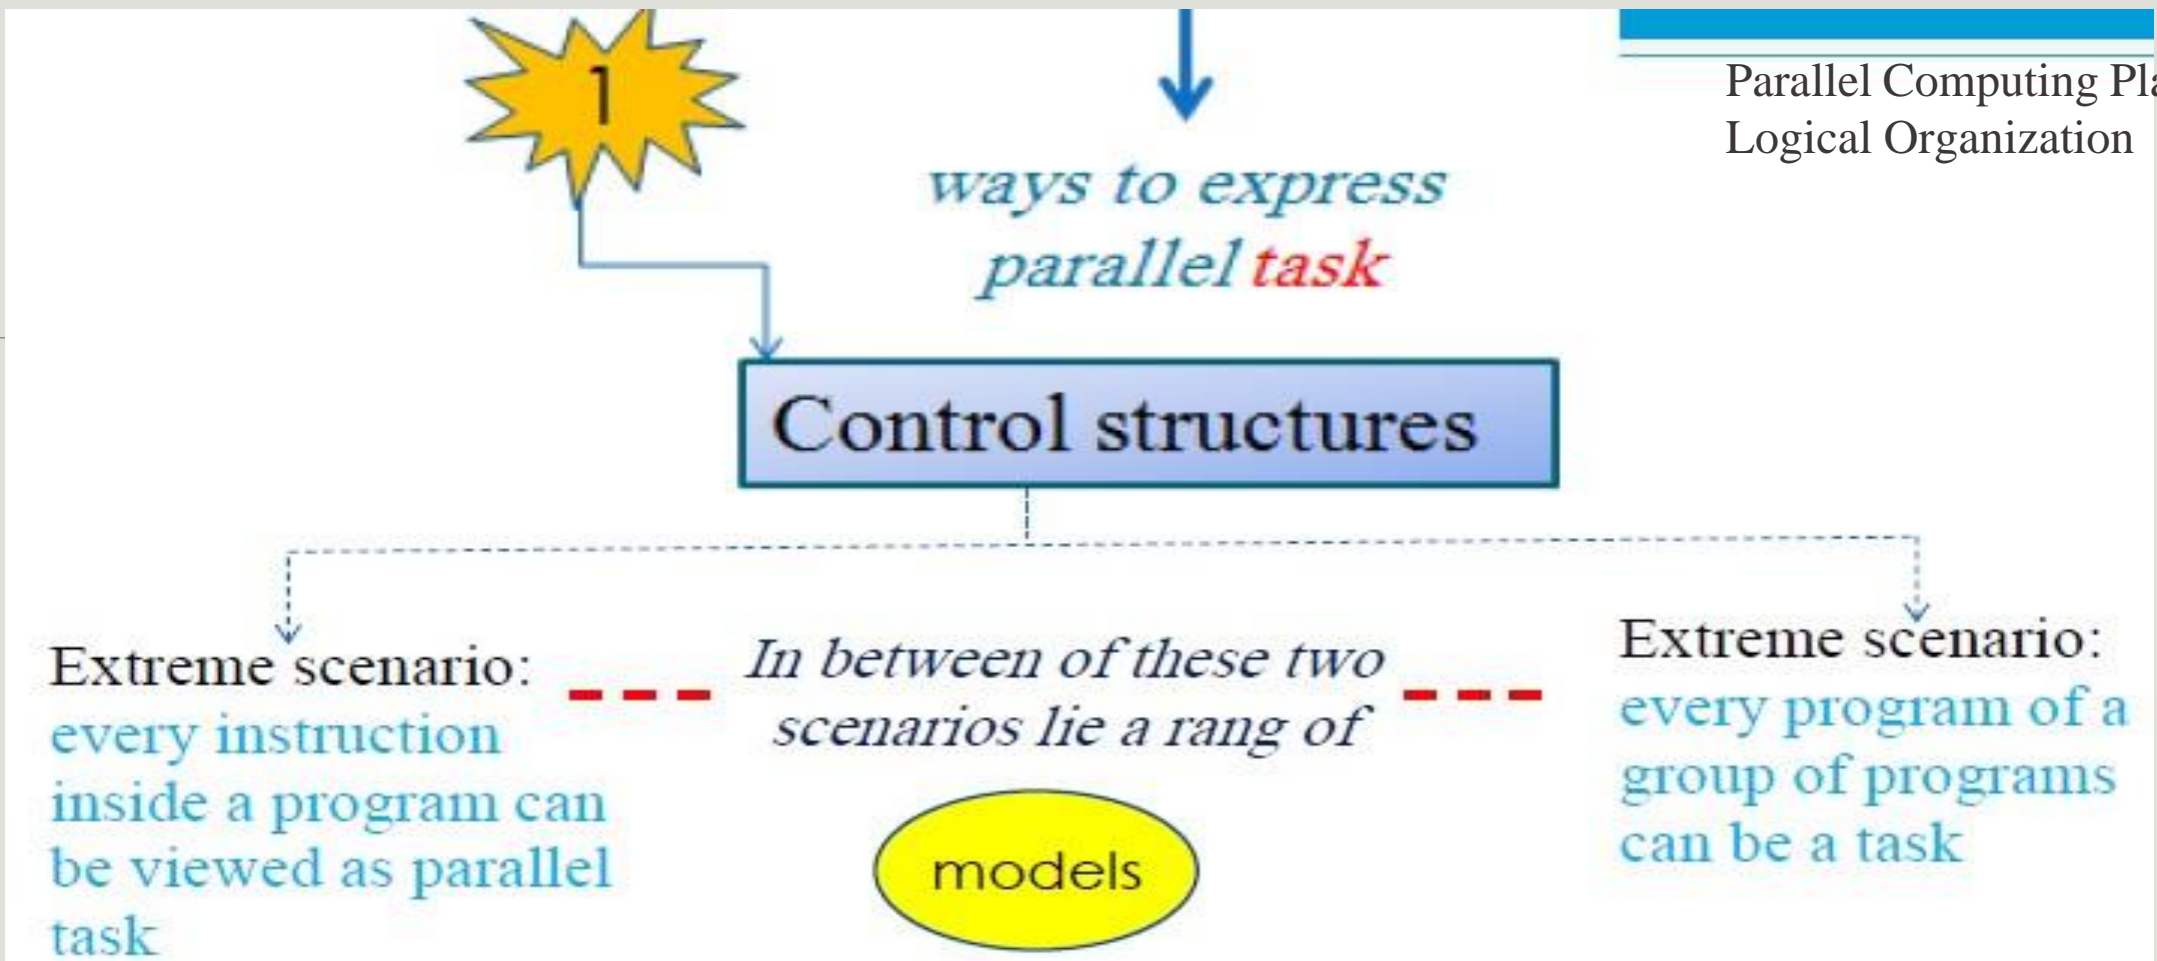
- ❖ The most efficient granularity is **dependent** on the **algorithm** and the **hardware environment** in which it runs.
- Use fine-grained granularity when the overhead of creating and managing tasks or threads is relatively low compared to the computational workload of each task.
- Fine-grained tasks are suitable when you need to maximize parallelism, and the tasks have limited data dependencies.
- Fine-grained granularity is appropriate when you want to keep all available CPU cores highly utilized.



Which is Best?

- ❖ The most efficient granularity is **dependent** on the **algorithm** and the **hardware environment** in which it runs.
- Use coarse-grained granularity when the overhead of managing a large number of fine-grained tasks becomes a significant performance bottleneck.
- Coarse-grained tasks are suitable when there are significant data dependencies or when managing fine-grained tasks becomes too complex.
- Coarse-grained granularity is appropriate when the computational workload of each task is substantial, and you want to reduce the overhead of task creation and management.





Task A logically **discrete section** of **computational work**. A task is typically a program or program-like set of instructions that is executed by a processor.



Models:

Flynn's Classical Taxonomy

- ❖ Flynn's classification scheme is based on the notion of a stream of information.
- ❖ Two types of information flow into a processor: **instructions** and **data**
- ❖ Each of these dimensions can have only one of two possible states: **Single** or **Multiple**.
- ❖ One of the more widely used classifications that classify parallel computers, use since 1966.



Flynn's Classical Taxonomy

The matrix below defines the 4 possible classifications according to Flynn:

models

SISD
Single Instruction, Single Data

SIMD
Single Instruction, Multiple Data

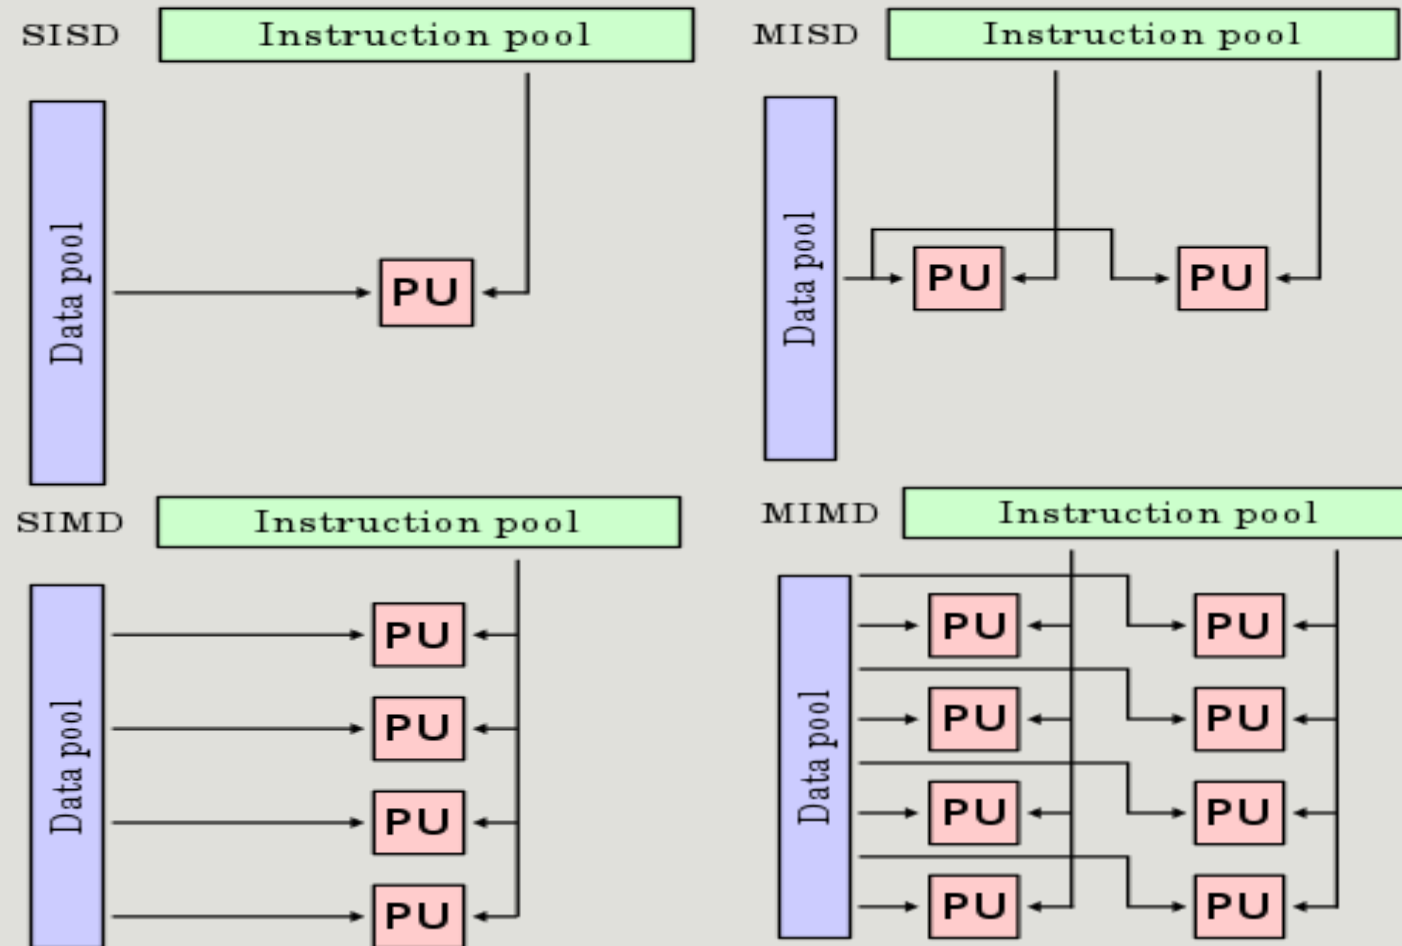
MISD
Multiple Instruction, Single Data

MIMD
Multiple Instruction, Multiple Data



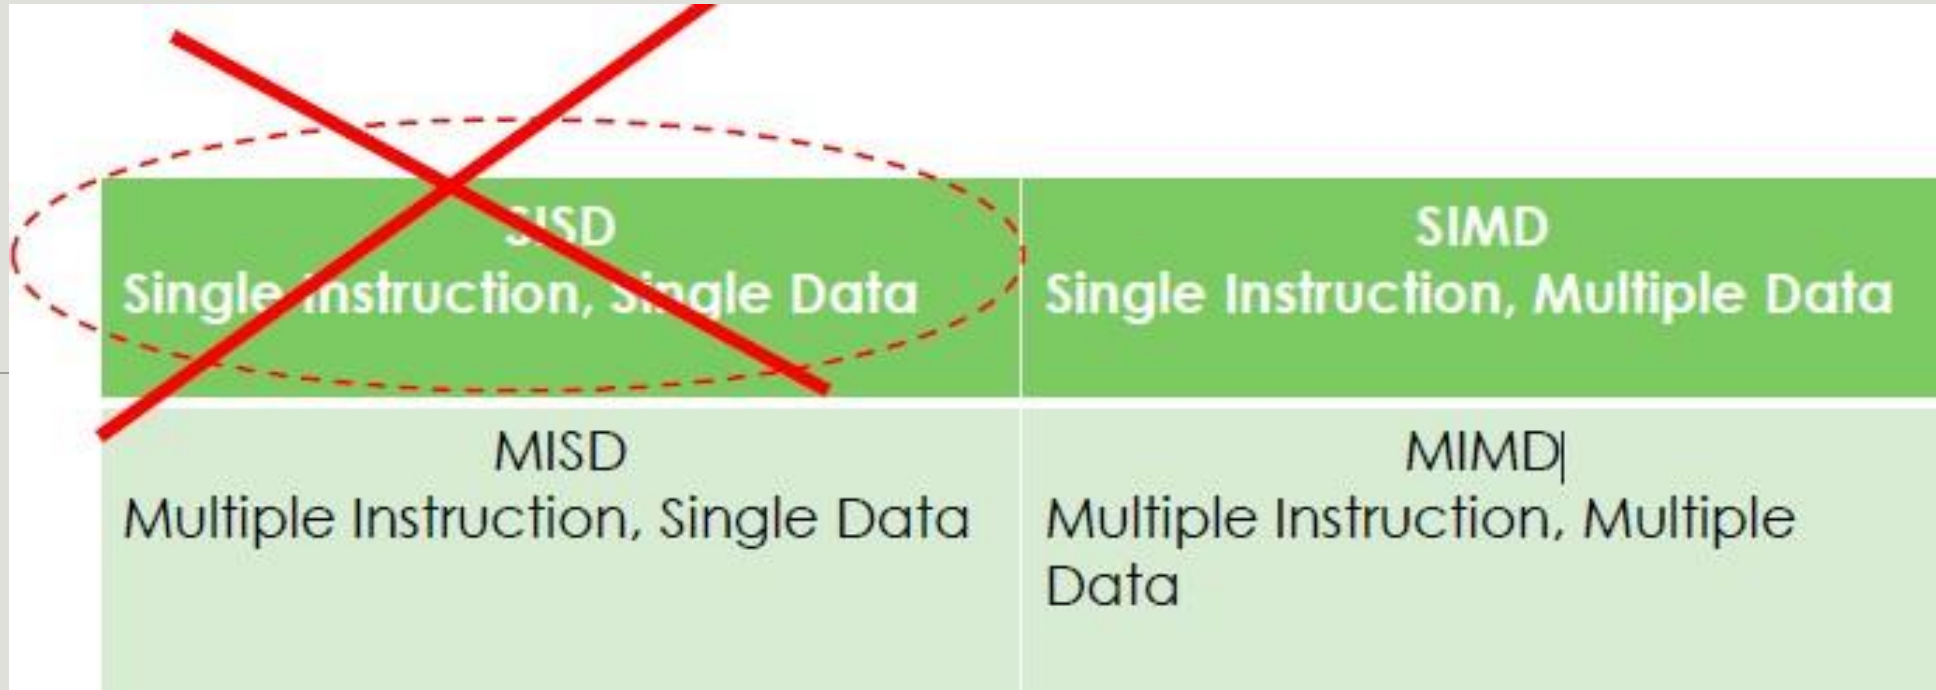
Models:

Flynn's Classical Taxonomy



Flynn's Classical Taxonomy

The matrix below defines the 4 possible classifications according to Flynn:

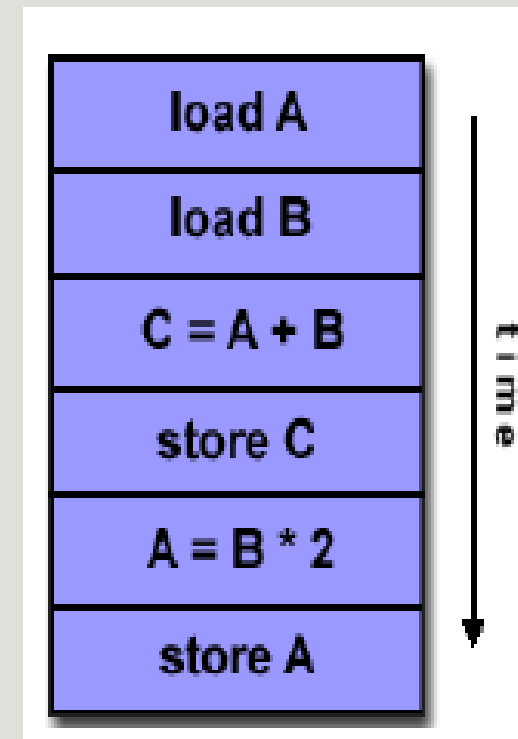


SISD Single Instruction, Single Data	SIMD Single Instruction, Multiple Data
MISD Multiple Instruction, Single Data	MIMD Multiple Instruction, Multiple Data



Single Instruction, Single Data (SISD):

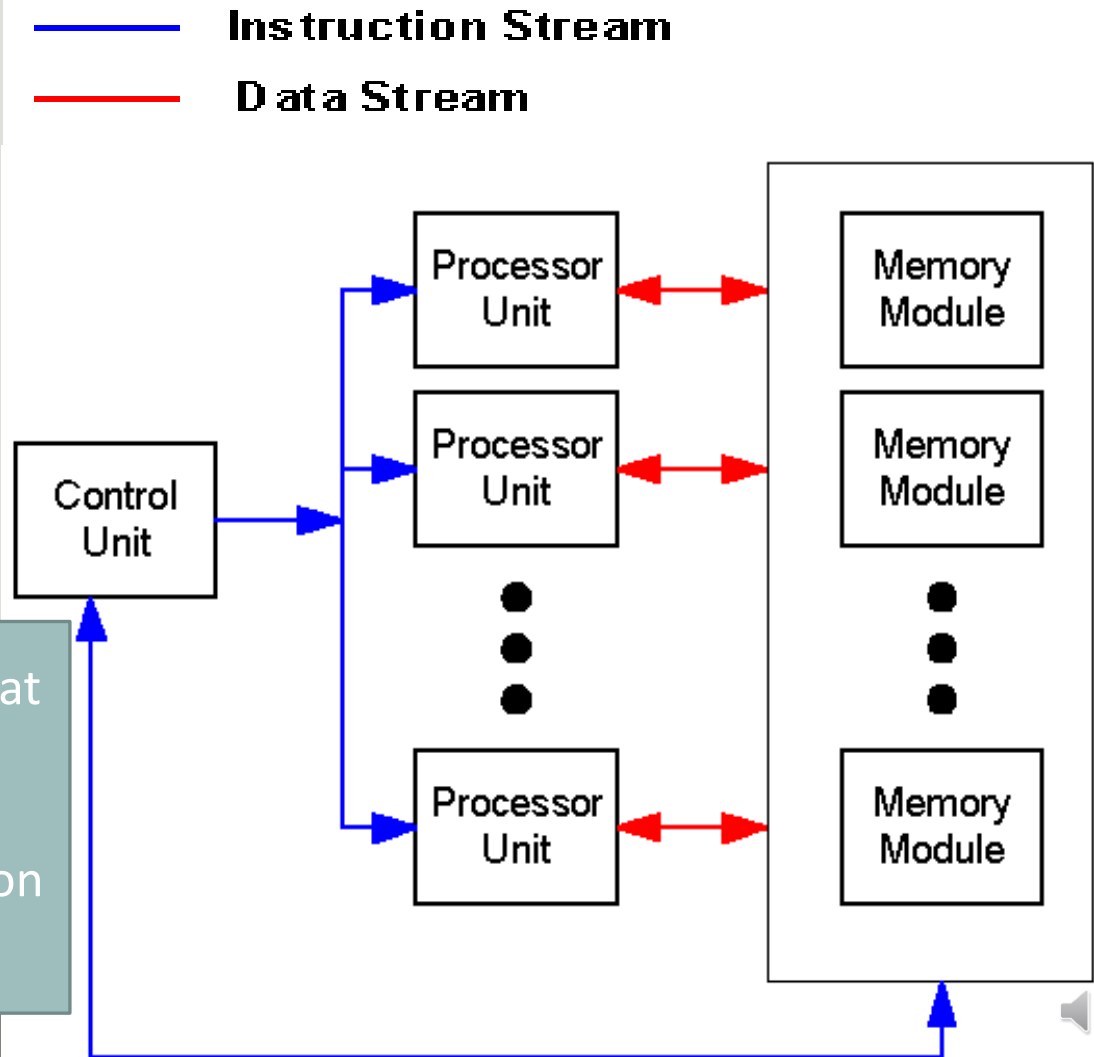
- ❖ A serial (non-parallel) computer
- ❖ Single instruction: only one instruction stream is being acted on by the CPU during any one clock cycle
- ❖ Single data: only one data stream is being used as input during any one clock cycle
- ❖ This is the oldest and even today, the most common type of computer
- ❖ Examples: older generation mainframes, minicomputers and workstations; most modern day PCs.




Single Instruction, Multiple Data (SIMD):

- ❖ Single instruction: All processing units execute the same instruction at any given clock cycle
- ❖ Multiple data: Each processing unit can operate on a different data element

A single control unit that dispatches the same instruction to various processors (that work on different data)




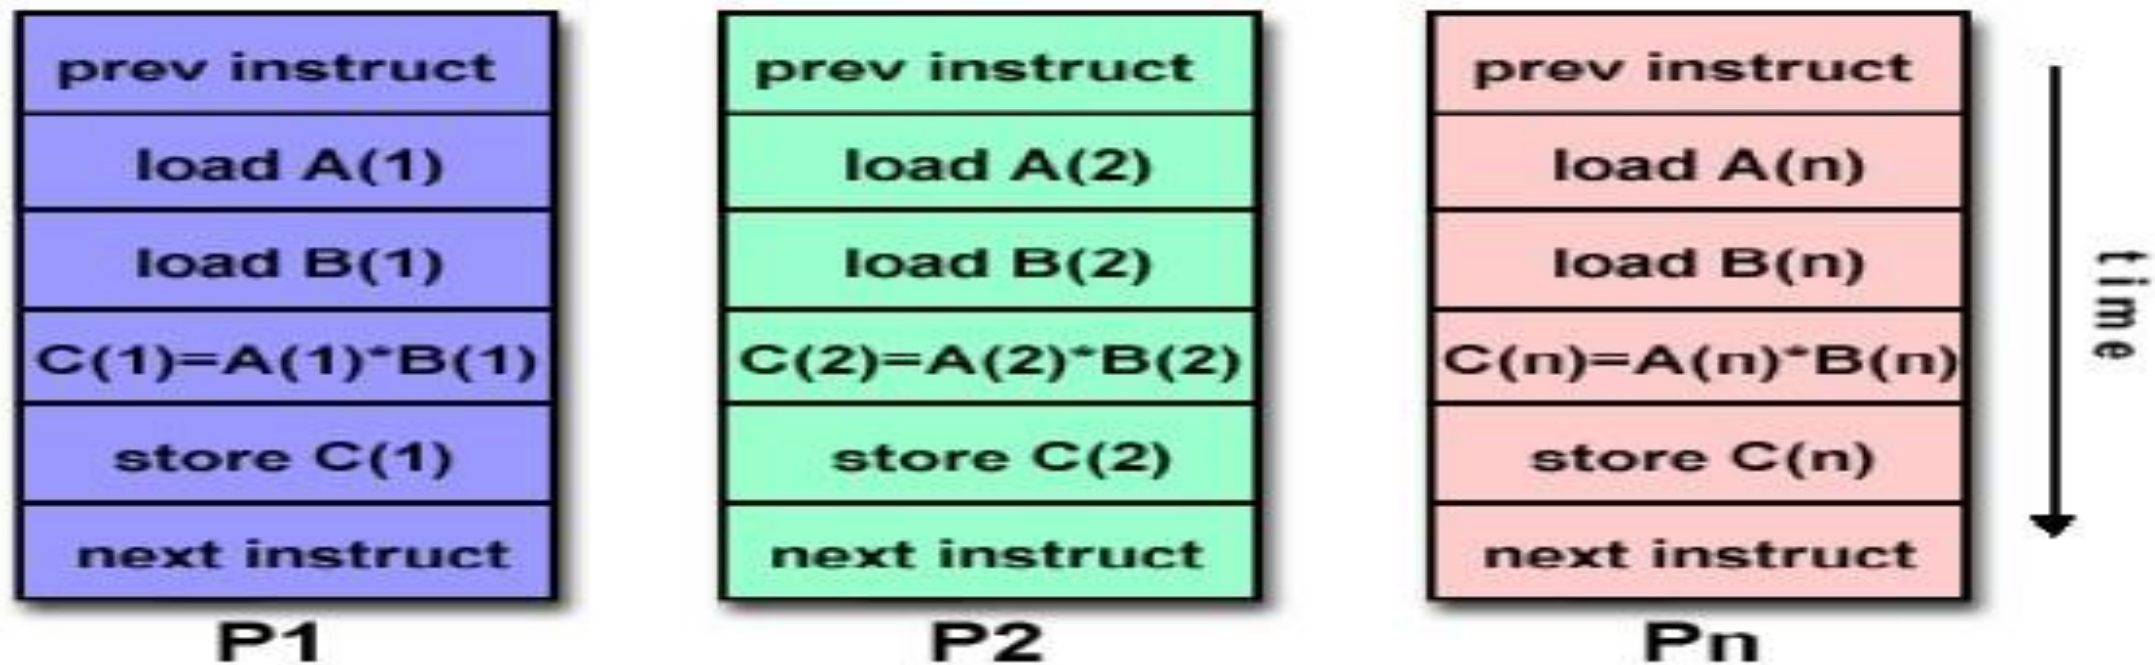


- ❖ Processor Arrays: ILLIAC IV, DAP Connection Machine CM-2, MasPar MP-1.
- ❖ Vector Pipelines: IBM 9000, Cray X-MP, Y-MP & C90, Fujitsu VP, NEC SX-2, Hitachi S820, ETA10
- ❖ Most modern computers, particularly those with graphics processor units (**GPUs**) employ SIMD instructions and execution units.
- ❖ Examples:

For (l = 0; i<1000; i++)

c[i] = a[i] * b[i];







```

if (B == 0)
    C = A;
else
    C = A/B;


```




Processor 0



Processor 1




Processor 2




Processor 3


Initial values




Processor 0



Processor 1




Processor 2




Processor 3


Step 1




Processor 0



Processor 1



Processor 2



Processor 3

Step 2

*One problem of
SIMD architectures*

Executing a conditional statement on an SIMD computer with four processors: (a) the conditional statement; (b) the execution of the statement in two steps.



Single Instruction, Multiple Data (SIMD):

❖ Conditional Operations in SIMD: Divergence leading to idle states

❖ Possible Solutions:

- SIMD friendly algorithm
- Prediction Hardware
- Branchless code (Ternary operator)
- Masked Operations (Intel AVX-512)
- SIMD Compiler Directives


Your Turn !!!

Guess what are the SIMD drawbacks??!!



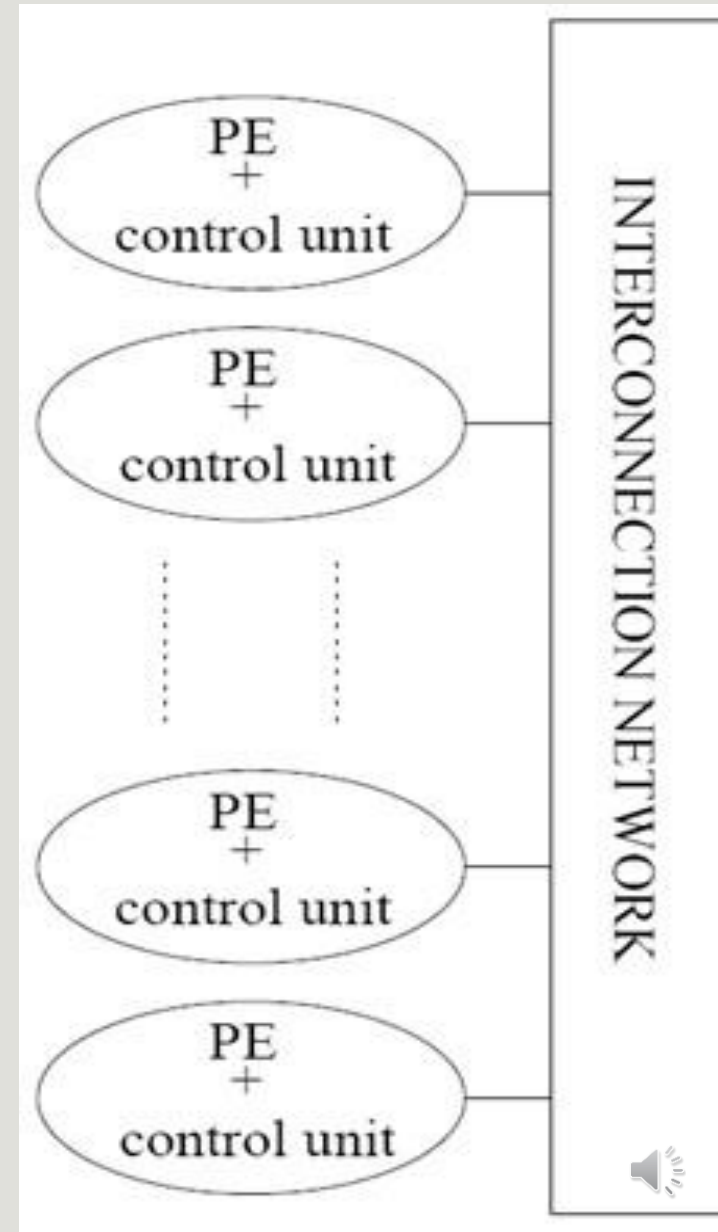
Multiple Instruction, Single Data (MISD):

- ❖ A single data stream is fed into multiple processing units.
- ❖ Each processing unit operates on the data independently via independent instruction streams.
- ❖ Few actual examples of this class of parallel computer have ever existed. One is the experimental Carnegie-Mellon C.mmp computer (1971).
- ❖ ex: Multiple cryptography algorithms attempting to crack a single coded message, fault tolerance.

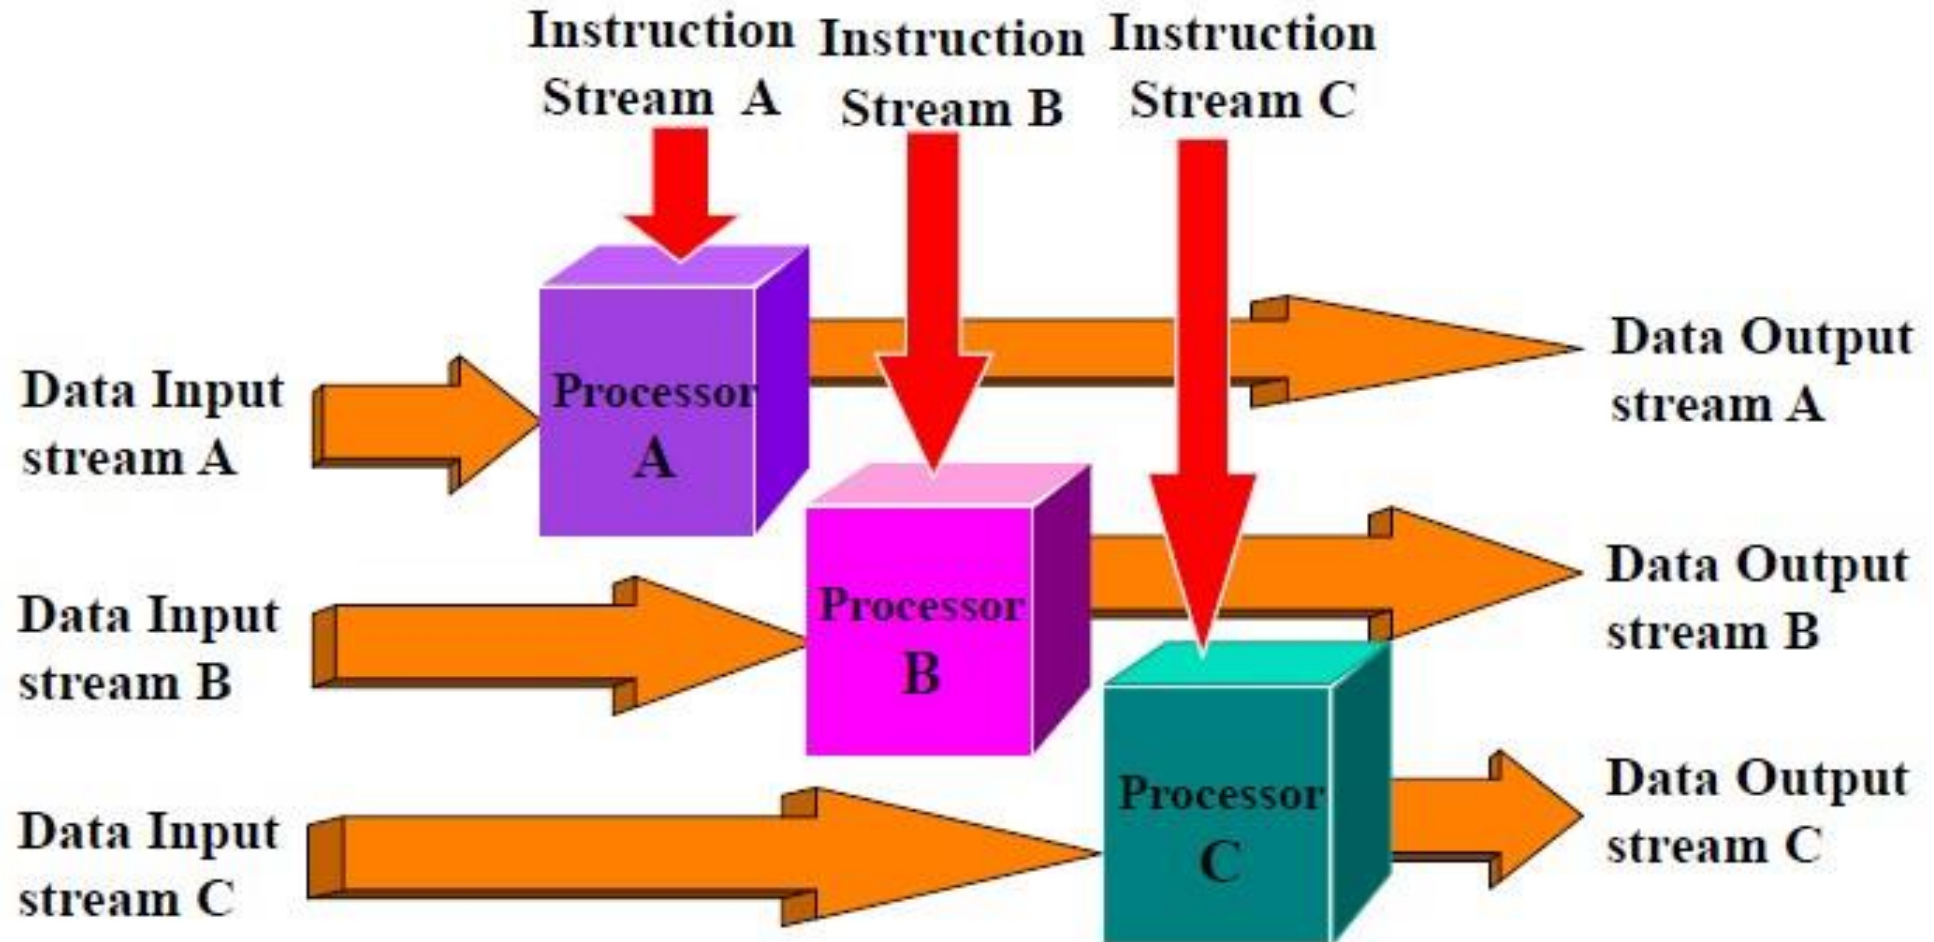


Multiple Instruction, Multiple Data (MIMD):

- ❖ Currently, the most common type of parallel computer. Most modern computers fall into this category.
- ❖ Multiple Instruction: every processor may be executing a different instruction stream.
- ❖ Multiple Data: every processor may be working with a different data stream.
- ❖ Examples: most current supercomputers, networked parallel computer **clusters** and "**grids**", **multi-processor SMP** computers, **multi-core** PCs.
- ❖ Note: many MIMD architectures also include SIMD execution sub-components




Multiple Instruction, Multiple Data (MIMD):



Multiple Instruction, Multiple Data (MIMD):

A simple variant of this model is

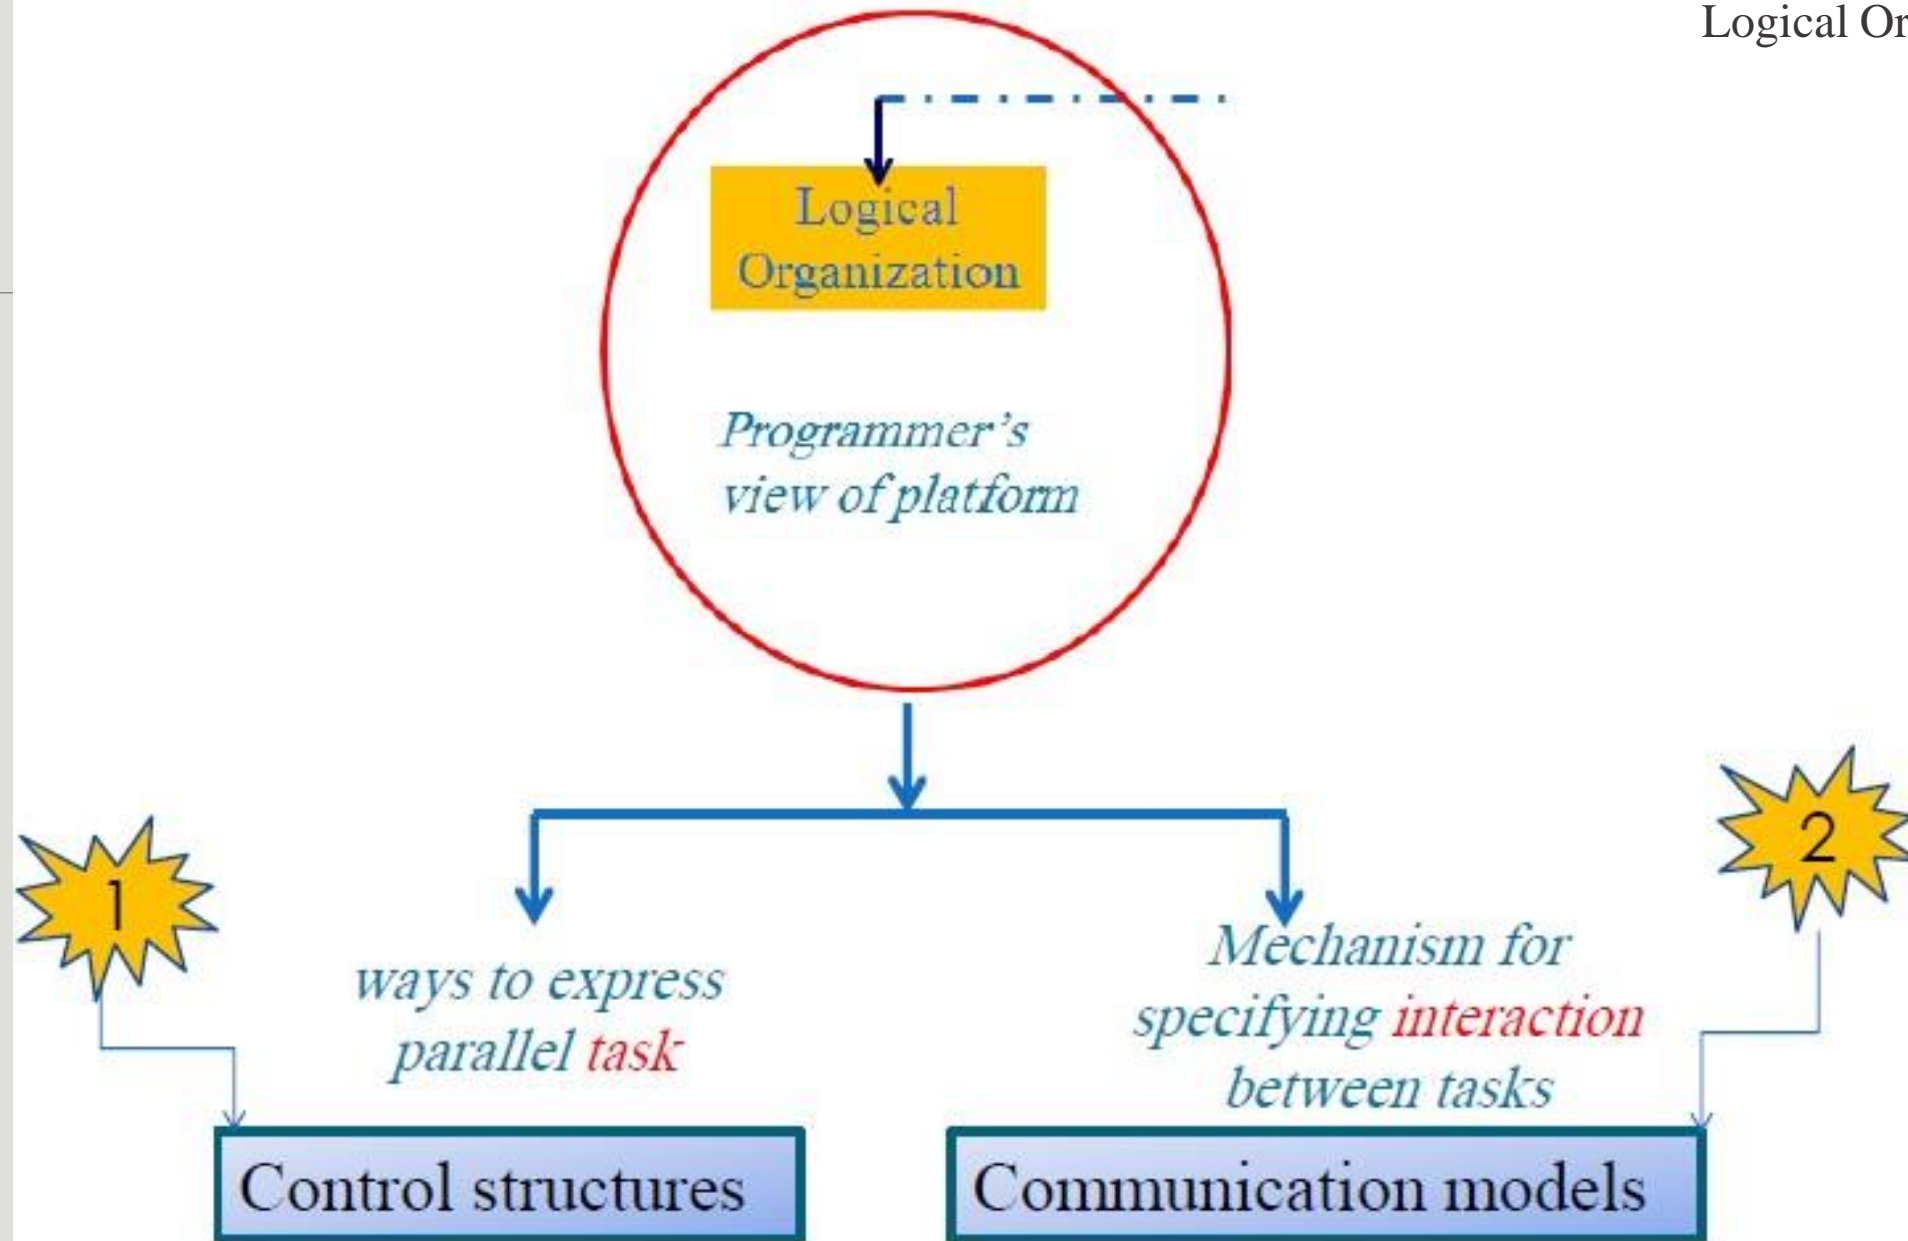
- ❖ Relies on multiple instance of the same program executing on different data
- ❖ Widely used by many parallel platforms and requires minimal architectural support
- ❖ Ex : Sun Ultra servers, multiprocessor PCs, workstation cluster & IBM SP



Your Turn

Compare between SIMD and MIMD



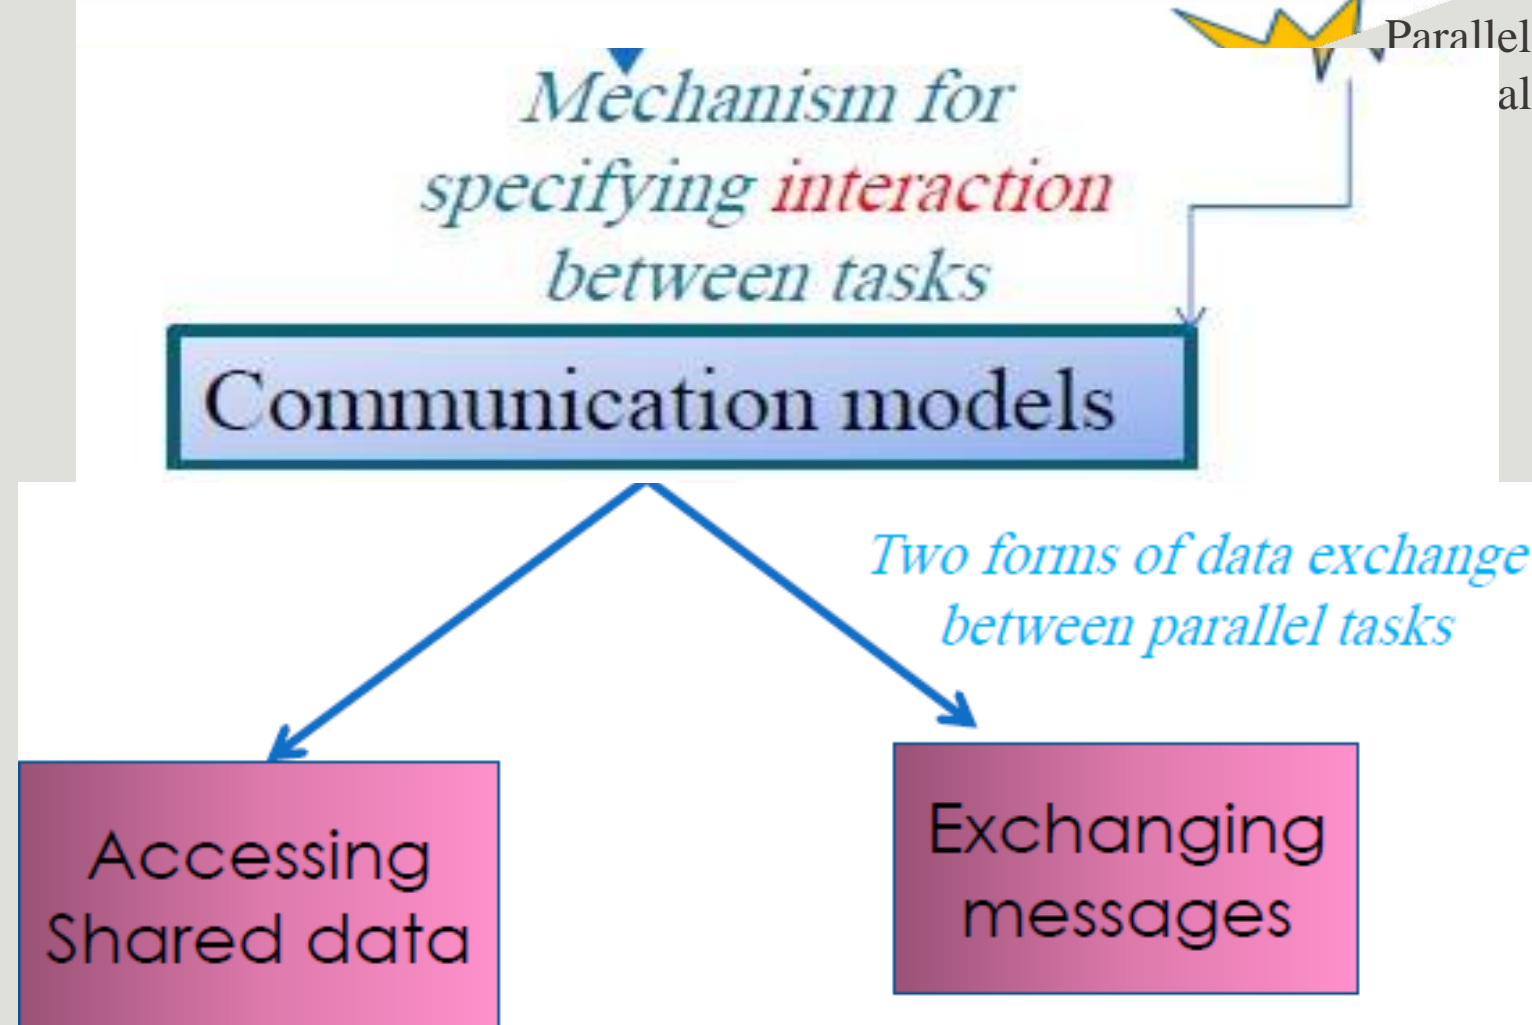


Communication Models

- ❖ Communication models in parallel computing define how processes or threads in a parallel system exchange data and information.
- ❖ These models specify the rules and mechanisms for communication among the components of a parallel system, which can include processes, threads, or even distributed nodes in a cluster or supercomputer.
- ❖ Different communication models are used depending on the architecture of the parallel system and the communication needs of the application.

Communication Models

- ❖ Shared Memory: Processes or threads communicate by reading and writing to a shared memory space..
- ❖ Distributed Memory: Each process has its own local memory, and communication between processes occurs by explicitly transferring data between them or message passing.
- ❖ Remote Procedure Call (RPC): Is a communication model where a process can invoke procedures or functions on a remote process as if they were local. It abstracts the communication details and provides a way for distributed processes to call functions on remote machines.
- ❖ Publish-Subscribe: Processes can subscribe to events or topics and publish messages. Subscribers receive messages published (e.g. MQTT).
- ❖ Shared Virtual Memory: An abstraction that allows distributed processes to access a common virtual address space [physically distributed](e.g. Trademark)



Platforms that provide a shared data space are called **shared-address-space machines** or **multiprocessors**

Platforms that support messaging are called **message passing platforms** or **multi-computers**.

